

Shivang Raikar

Full-stack and AI engineer

San Francisco, CA | 857-366-3829 | shivangraikar@gmail.com | [LinkedIn](#) | [GitHub](#) | [Medium](#) | [Scholar](#)

SUMMARY

Full-stack and AI engineer who takes products from an empty repo to real users. As a core member of the founding team at Ohm, I built our main mobile application, client-facing APIs for data exchange, operational web application to schedule, monitor, and track all product services across customers. My technology strengths include application development, distributed systems, database management, CI/CD, logging, and deployment for real-world applications. Strong research background in machine learning and reinforcement learning; adopted AI workflows in daily coding and non-coding tasks to utilize models and tools that extend the problem-solving and solution-building process.

PROFESSIONAL EXPERIENCE

OHM - Founding Software Engineer

San Francisco, CA | Jan 2025 - Present

- Architected and shipped a multi-agent AI system for natural-language querying of operational data (payments, invoices, service history, client and customer records), leveraging LLM orchestration patterns, tool calling, and knowledge extraction with fuzzy search logic for accurate responses to our customers, clients, and admins.
- Lead architect on "OhmControl", an operations platform using Next.js, React, and Django, powered by our AI agent Ohmie that handles operations for field service, reducing lookup time by implementing TanStack Query-based caching with context management to accelerate internal decision-making, reducing response latency from 400ms to <10ms.
- Shipped a React Native mobile application on the Play Store and App Store used by field technicians to perform on-site inspections and servicing on EV chargers and HVAC systems, capturing real-time diagnostic data that feeds directly into customer-facing reporting dashboards.
- Integrated Ohmie across applications and extended its capabilities: worktype builder, service diagnosis, scheduling next visits, document parser for knowledge on service types, summarizing logistics for clients.

STEAM WORKS STUDIO - Software Engineer

Downingtown, PA | Aug 2024 - Jan 2025

- Led a full-stack RAG-powered document intelligence system using Python, Django, and LangChain, enabling natural-language querying across curriculum materials, student portfolios, and assessment data for K-12 education institutions.
- Built a PDF and document parsing pipeline processing 1,000+ pages of educational content with chunking, embedding generation (OpenAI), and vector storage to power semantic search across chatbot systems using Reducto and LlamaParse.
- Expanded the platform to serve clients in insurance, clinical, and legal domains, supporting 200+ daily active users with sub-200ms response times and measurable improvements in conversational accuracy.

MPSYCH LAB, UMASS BOSTON - Research Software Engineer

Boston, MA | Aug 2023 - Feb 2024

- Leveraged University of Chicago datasets for cancer patient analysis; built an interactive visualization tool for treatment outcome leaderboards, enabling clinicians to compare protocol efficacy across 5,000+ patient records.
- Optimized 12+ API scripts and serializers, enabling real-time data capture and transfer between disparate systems, improving data accuracy from raw AWS S3 buckets by 50%.
- Engineered comprehensive data handling and migration strategies for 10,000+ records, leveraging MySQL and Pandas to enhance data integrity and accessibility for cross-functional engineering teams.

CRANE PAYMENT INNOVATIONS - Software Engineer

Malvern, PA | Jun 2022 - Dec 2022

- Developed and automated a transaction-based verification system for validating incoming and outgoing cash and coin flows across hardware devices deployed in 10,000+ casinos and vending machines nationwide.
- Designed 4 POCs to validate core business logic, particularly ensuring test reliability on admin portals, internal functions, and modules that reduced manual cash audit time by 65%, processing \$2M+ in daily transactions with 99.9% accuracy.
- Eliminated 40% of manual testing by automating JUnit scripts with AWS Lambda workflows, improved querying speed on Elasticsearch by 25%, optimized indexing, joins, and prefetch.

SKILLS & CERTIFICATIONS

- **Languages:** Python, TypeScript, JavaScript, SQL, Go, Java, C++
- **Frontend & Mobile:** React, Next.js, Redux, Zustand, Expo, ReactNative, Streamlit
- **Tools and IDE:** Git, PyUnit, Selenium, JUnit, CI/CD, Postman, Celery, Claude Code, Docker
- **Databases/Cloud:** Supabase, Firebase, MongoDB, PostgreSQL, DynamoDB, Redis, AWS, GCP
- **AI/ML:** LLMs, PyTorch, agentic AI, LangChain, RAG, vector database, model deployment, reinforcement learning
- **Certifications:** [AWS AI Practitioner](#), [HackerRank Software Engineer](#), [Oracle Cloud Generative AI Professional](#)

EDUCATION

Master of Science in Computer Science - **University of Massachusetts**

Boston, MA | May 2023

Bachelor of Engineering in Computer Engineering - **University of Mumbai**

Mumbai, India | May 2021

ACHIEVEMENTS

- Co-authored "[Twitter Data Mining for Targeted Marketing](#)", 2020 Second International Conference on Inventive Research in Computing Applications (ICIRCA), Coimbatore, India, under the Machine Learning domain.
- Software Engineer Fellow @Headstarter AI, projects: shivangraikar.github.io (Developed 4+ full-stack SaaS products using TypeScript, NextJS, MaterialUI, Stripe API, LLMs, Pinecone Vector DB, leading 3 developers.)
- Judged a hackathon at AWS Builder Loft, San Francisco; shortlisted 5 teams out of 40+ projects reviewed in the agentic AI and Machine learning space, sponsored by HydraDB, Photon, RocketRide, and more.
- Won and participated in 10+ hackathons, demonstrating rapid prototyping and cross-functional collaboration [link](#).